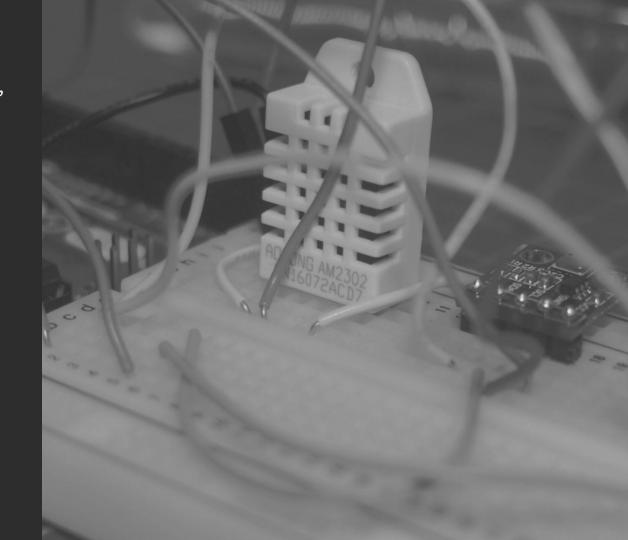
Local Weather Visualisation with IoT, Kafka and CrateDB







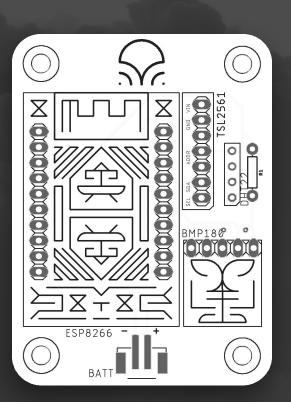
WHY BUILD A WEATHER STATION?

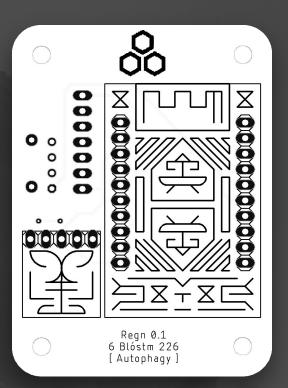


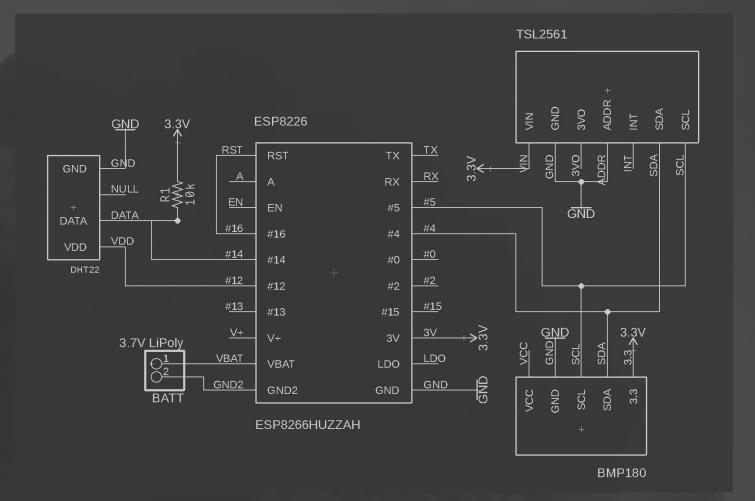


Temperature/Humidity
Barometric Pressure
Luminosity
Microcontroller

DHT22 BMP180 TSL2561 ESP8266







ACT 1 A SIMPLE PROTOTYPE

```
// Libraries
#include <DHT.h>
#include <Adafruit_Sensor.h>
#include <Adafruit_TSL2561_U.h>
#include <Adafruit_BMP085_U.h>
#include <ESP8266WiFi.h>
#include <ArduinoJson.h>
// Constants
#define DHTPIN 12 // DHT22 connected to pin 12
#define DHTTYPE DHT22
DHT dht(DHTPIN, DHTTYPE);
Adafruit_TSL2561_Unified tsl = Adafruit_TSL2561_Unified(TSL2561_ADDR_LOW, 12345);
Adafruit_BMP085_Unified bmp = Adafruit_BMP085_Unified(10085);
```

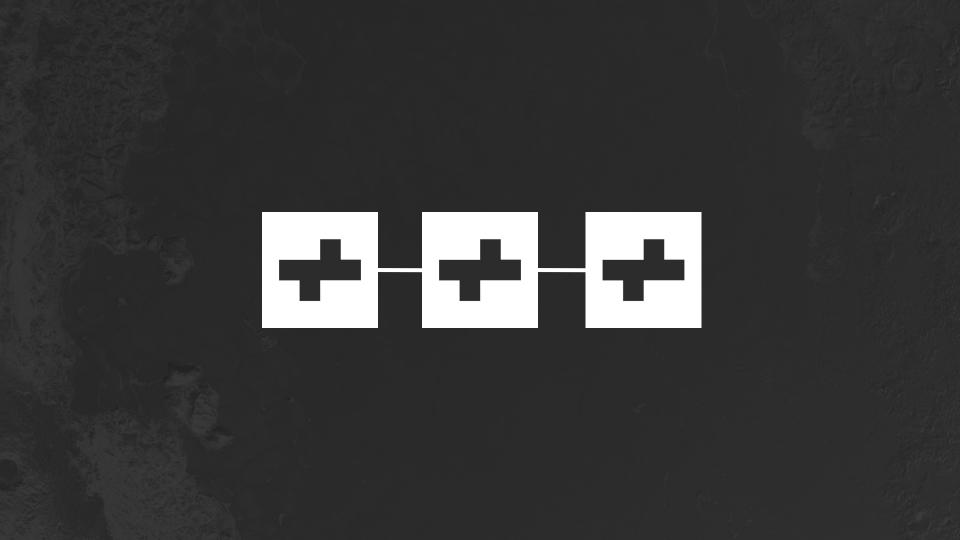
```
dht.begin();
tsl.enableAutoRange(true);
tsl.setIntegrationTime(TSL2561_INTEGRATIONTIME_13MS);
bmp.begin();
WiFi.begin(ssid, password);
while (WiFi.status() != WL_CONNECTED) {
    delay(500);
}
```

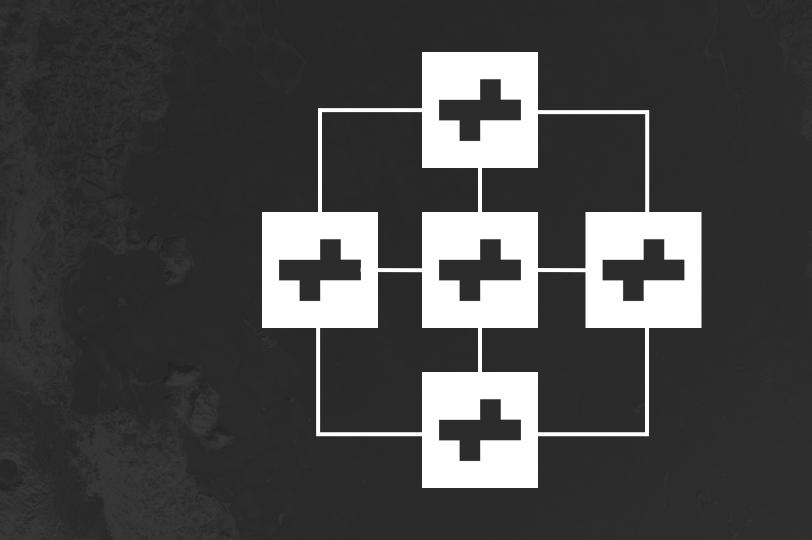
```
temperature = dht.readTemperature();
humidity = dht.readHumidity();
sensors_event_t bmp_event;
bmp.getEvent(&bmp_event);
pressure = bmp_event.pressure;
sensors_event_t tsl_event;
tsl.getEvent(&tsl_event);
lux = tsl_event.light;
```

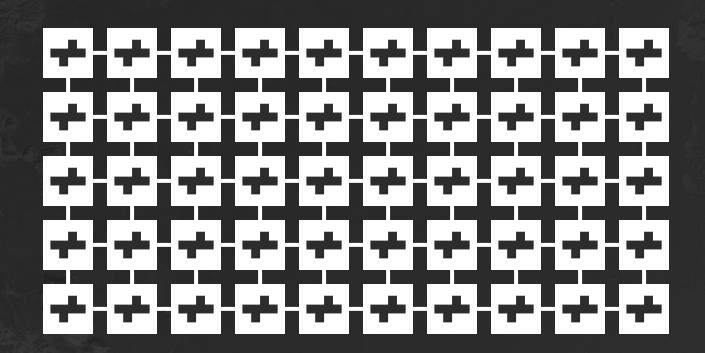
```
StaticJsonBuffer<256> JSONbuffer;
JsonObject& JSONencoder = JSONbuffer.createObject();
JSONencoder["temperature"] = temperature;
JSONencoder["humidity"] = humidity;
JSONencoder["pressure"] = pressure;
JSONencoder["luminosity"] = lux;
char JSONmessageBuffer[256];
JSONencoder.prettyPrintTo(JSONmessageBuffer, sizeof(JSONmessageBuffer));
```

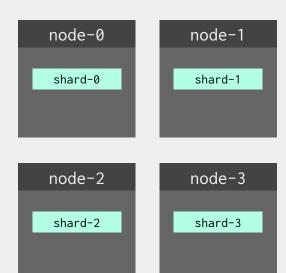
```
int code = -1;
HTTPClient http;
while(code != 201) {
  http.begin(host, fingerprint);
  http.addHeader("Content-Type", "application/json");
  code = http.POST(JSONmessageBuffer);
  http.end();
  delay(1000);
}
delay(sensor_interval);
```



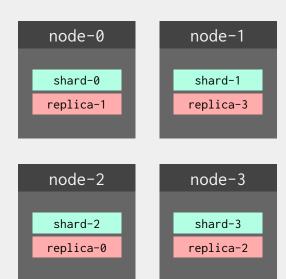




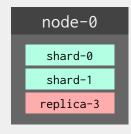


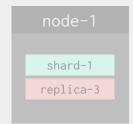


```
CREATE TABLE IF NOT EXISTS
"doc"."sensordata" (
    "timestamp" TIMESTAMP,
    "value" FLOAT
)
CLUSTERED INTO 4 SHARDS
```



```
CREATE TABLE IF NOT EXISTS
"doc"."sensordata" (
    "timestamp" TIMESTAMP,
    "value" FLOAT
)
CLUSTERED INTO 4 SHARDS
WITH (number_of_replicas = 1)
```





```
CREATE TABLE IF NOT EXISTS
"doc"."sensordata" (
    "timestamp" TIMESTAMP,
    "value" FLOAT
)
CLUSTERED INTO 4 SHARDS
WITH (number_of_replicas = 1)
```

```
node-2

shard-2

replica-0

replica-1
```

```
node-3
shard-3
replica-2
```

value	timestamp
12.4	1532080800000
68.2	1532080850000
107	1532080900000
12.2	1526810400000
69.0	1526810423000
110	1526810434000
11.9	1524219010000
69.1	1524219140000
109	1524219230000

```
CREATE TABLE IF NOT EXISTS
"doc"."sensordata" (
    "timestamp" TIMESTAMP,
    "value" FLOAT
)
CLUSTERED INTO 4 SHARDS
WITH (number_of_replicas = 1)
```

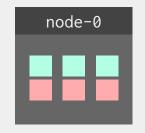
value	timestamp	month
12.4	1532080800000	1532080800000
68.2	1532080850000	1532080800000
107	1532080900000	1532080800000
12.2	1526810400000	1526810400000
69.0	1526810423000	1526810400000
110	1526810434000	1526810400000
11.9	1524219010000	1524218400000
69.1	1524219140000	1524218400000
109	1524219230000	1524218400000

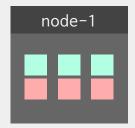
```
CREATE TABLE IF NOT EXISTS
"doc"."sensordata" (
    "timestamp" TIMESTAMP,
    "value" FLOAT,
    "month" TIMESTAMP GENERATED ALWAYS
    AS date_trunc('month', "timestamp")
)
CLUSTERED INTO 4 SHARDS
WITH (number_of_replicas = 1)
```

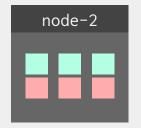
value	timestamp	month
12.4	1532080800000	1532080800000
68.2	1532080850000	1532080800000
107	1532080900000	1532080800000
12.2	1526810400000	1526810400000
69.0	1526810423000	1526810400000
110	1526810434000	1526810400000
11.9	1524219010000	1524218400000
69.1	1524219140000	1524218400000
109	1524219230000	1524218400000

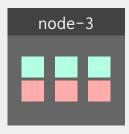
```
CREATE TABLE IF NOT EXISTS
"doc"."sensordata" (
    "timestamp" TIMESTAMP,
    "value" FLOAT,
    "month" TIMESTAMP GENERATED ALWAYS
    AS date_trunc('hour', "timestamp")
)
CLUSTERED INTO 4 SHARDS
```

PARTITIONED BY ("month")
WITH (number_of_replicas = 1)









CREATE TABLE IF NOT EXISTS

```
"doc"."sensordata" (
    "timestamp" TIMESTAMP,
    "value" FLOAT,
    "month" TIMESTAMP GENERATED ALWAYS
    AS date_trunc('hour', "timestamp")
)
CLUSTERED INTO 4 SHARDS
PARTITIONED BY ("month")
WITH (number_of_replicas = 1)
```





?





```
@main.route("/api/insert", methods=["POST"])
@validate_schema(insert_sensor_schema)

def insert():
    data = json.loads(request.data)
    timestamp = datetime.utcnow()
    db_data = SensorData(
        timestamp=timestamp,
```

temperature=data["temperature"],

humidity=data["humidity"],
pressure=data["pressure"],
luminosity=data["luminosity"],

db.session.add(db_data)

db.session.commit()

return "", 201

```
@main.route("/api/latest", methods=["GET"])

def latest():
    latest_results = COORDINATES

    latest = db.session.query(SensorData).order_by(desc("timestamp")).first()
    if (latest is not None):
```

latest_results.update(latest.dict())

return jsonify(latest_results)

```
mika@autophagy > curl localhost:8080/api/latest
```

```
"latitude": 52.489,
"longitude": 13.354,
```

"temperature": 11.2,

"humidity": 91,

```
"pressure": 1012,
```

```
"lux": 107
```

```
@main.route("/api/<string:sensor>/since/<int:ts>", methods=["GET"])

def sensorReadingsSinceTimestamp(sensor, ts):
    if sensor not in VALID_SENSOR_TYPES:
        abort(404)

    dt = datetime.utcfromtimestamp(ts / 1000)
    values = db.session.query(getattr(SensorData, sensor), SensorData.timestamp).filter(
```

SensorData.timestamp > dt

asc("timestamp")

).order_by(

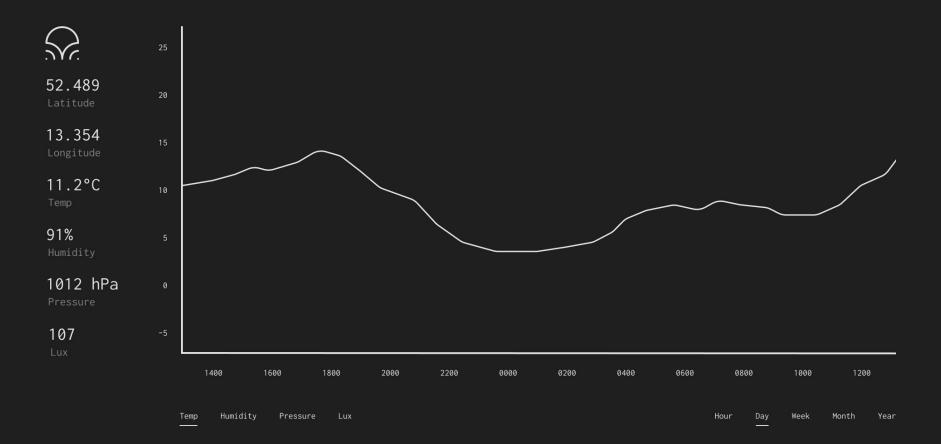
return jsonify(values)

).all()

	{ "timestamp": 1531920181000, "value": 12.2 },
]	

mika@autophagy > curl localhost:8080/api/temperature/since/1531919792197

{ "timestamp": 1531920179000, "value": 12.0 }, { "timestamp": 1531920180000, "value": 12.4 },



INTERLUDE



SCALABILITY

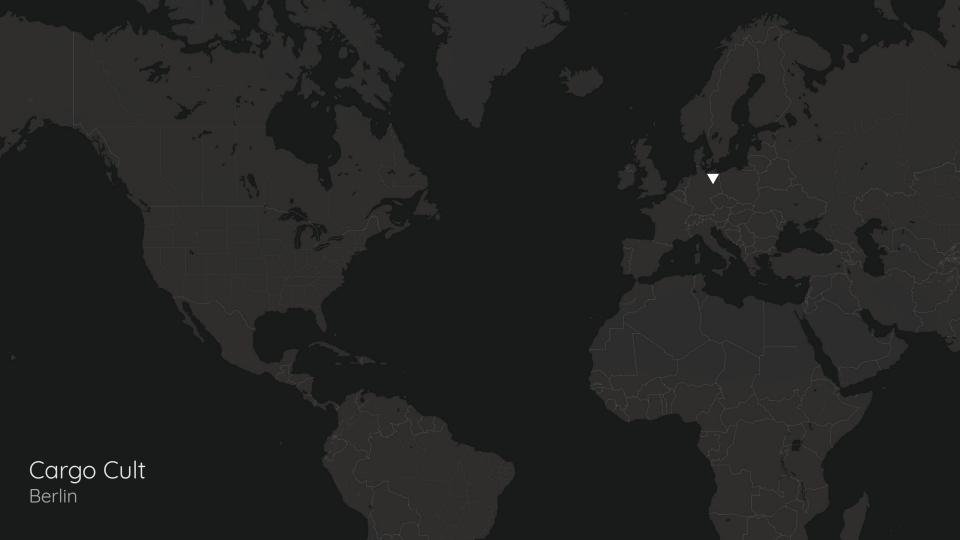
FRIENDS WANTED A REGN STATION FOR THEIR OWN PURPOSES

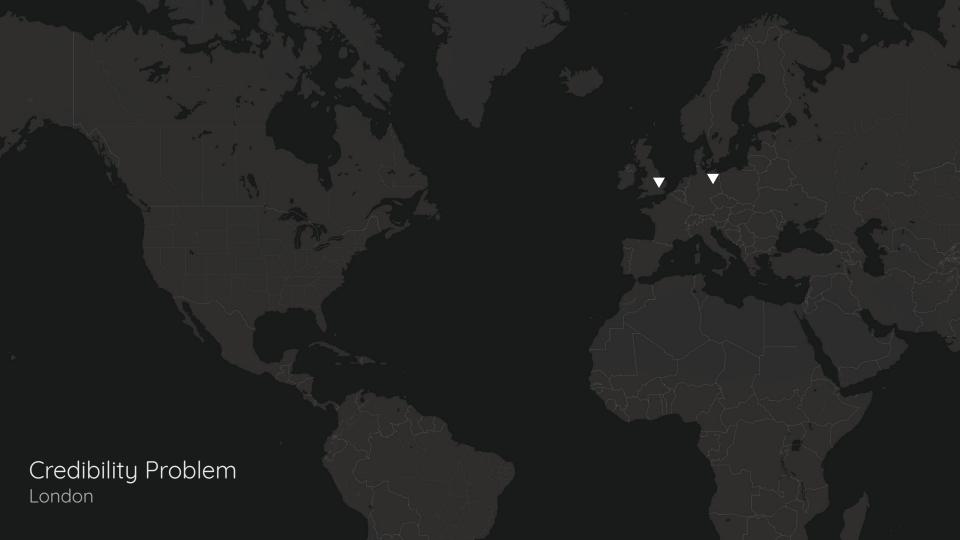
ATOMISED SENSORS or A SENSOR NETWORK

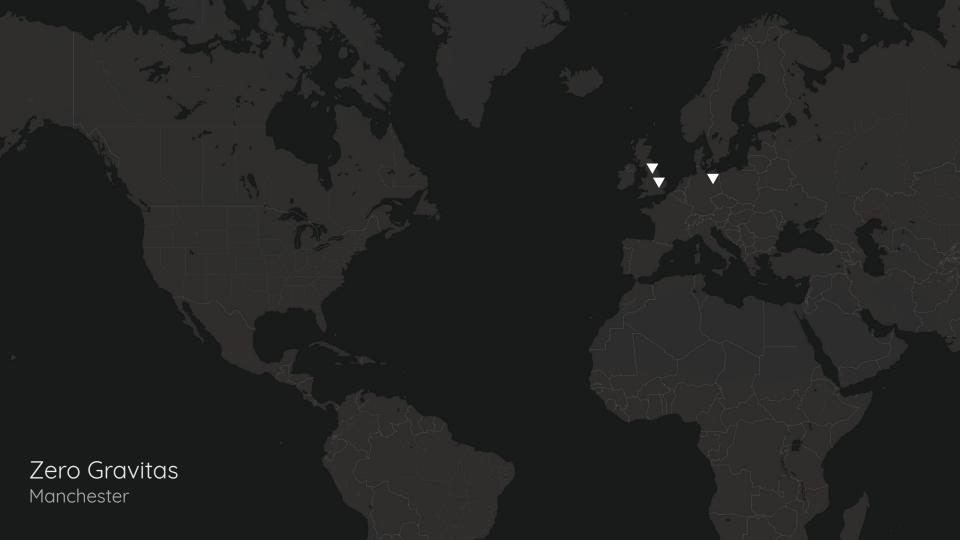
THE GANG MAKES A WEATHER

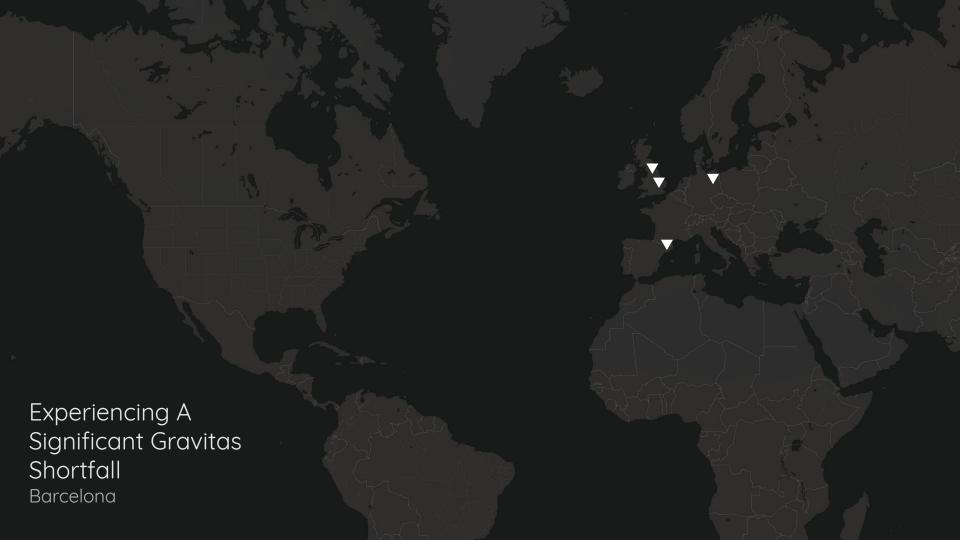
ACT 2

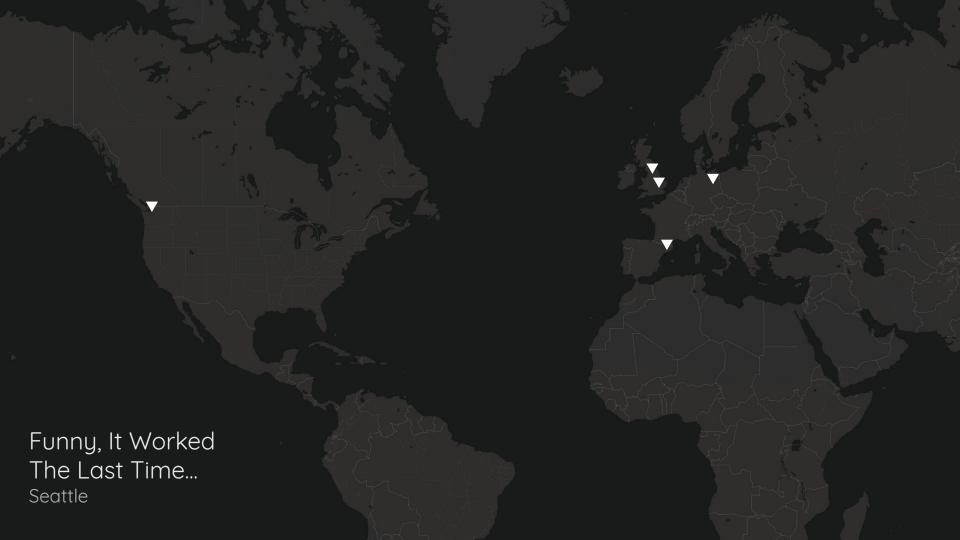
STATION NETWORK



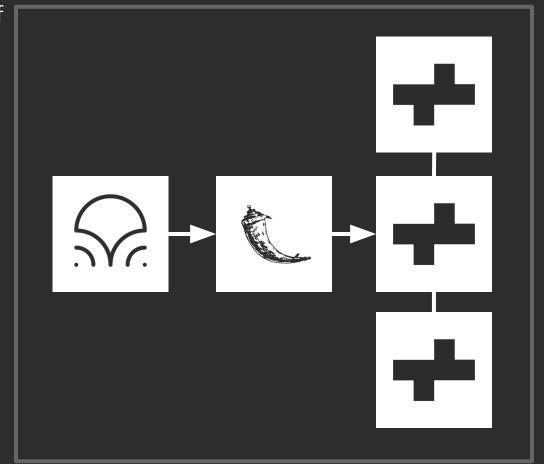




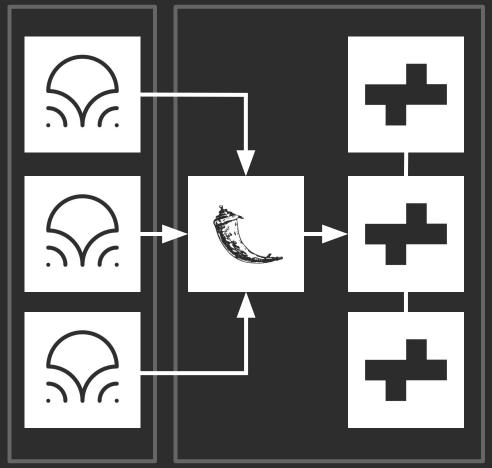




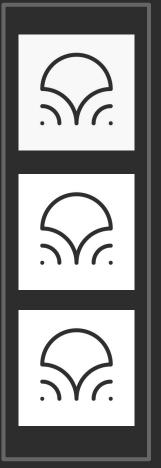
My Stuff

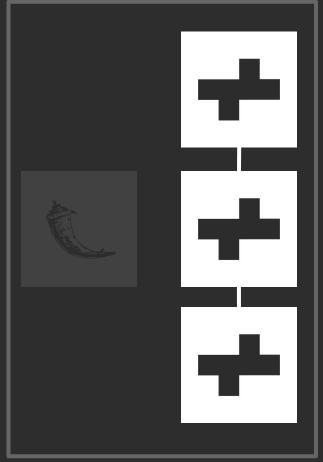


Not My Stuff

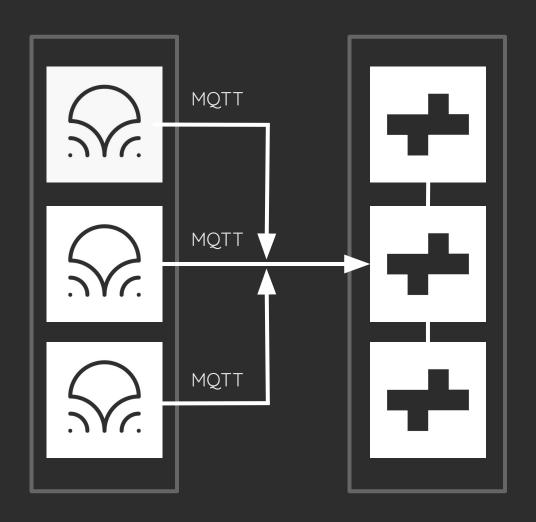


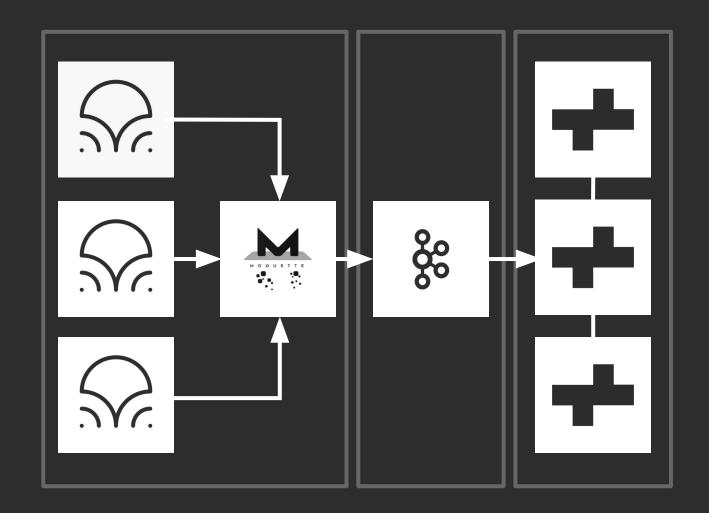
My Stuff? Our Stuff? Not My Stuff





My Stuff? Our Stuff?





```
// Libraries
#define MQTT_MAX_PACKET_SIZE 256
#include <PubSubClient.h>

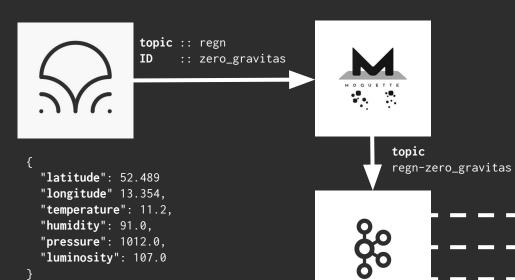
// Constants
WiFiClient espClient;
PubSubClient client(espClient);
```

```
void callback(char* topic, byte* payload, unsigned int length) {
    Serial.print("Received Message :: ");
    Serial.print(topic);
    for (int i = 0; i < length; i++) {
        Serial.print((char)payload[i]);
    }
    Serial.println();
}</pre>
```

```
client.setServer(mqtt_server, 1883);
client.setCallback(callback);

if (client.connect("zero_gravitas")) {
   client.subscribe("regn/zero_gravitas");
}
```

```
StaticJsonBuffer<256> JSONbuffer;
JsonObject& JSONencoder = JSONbuffer.createObject();
JSONencoder["latitude"] = sensor_latitude;
JSONencoder["longitude"] = sensor_longitude;
JSONencoder["temperature"] = temperature;
JSONencoder["humidity"] = humidity;
JSONencoder["pressure"] = pressure;
JSONencoder["luminosity"] = lux;
char JSONmessageBuffer[256];
JSONencoder.prettyPrintTo(JSONmessageBuffer, sizeof(JSONmessageBuffer));
client.publish("regn", data);
delay(sensor_interval);
```



CONSUMER

CONSUMER

CONSUMER

CONSUMER

CONCLUSION LESSONS // FUTURE WORK

ESP8266 IS A GREAT INTRO TO IOT PROJECTS

TIGHTLY COUPLED IOT DATA PIPELINES ARE KINDA BAD

CRATEDB & KAFKA ARE PRETTY AWESOME

FUTURE LEVERAGE MQTT MORE BROADEN CRATEDB INGESTION EXPERIMENT WITH OUR DATA!



Twitter Github Email @autophagyDevautophagy mika@crate.io